



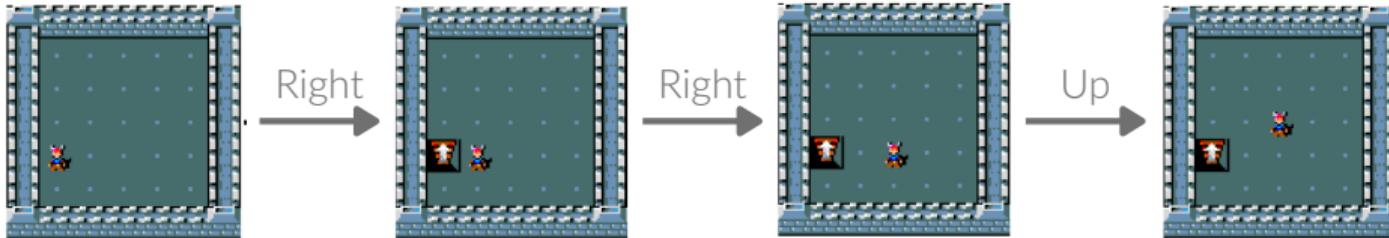
WILLEBERG AI
AUTONOMOUS SYSTEMS
AND SOFTWARE PROGRAM



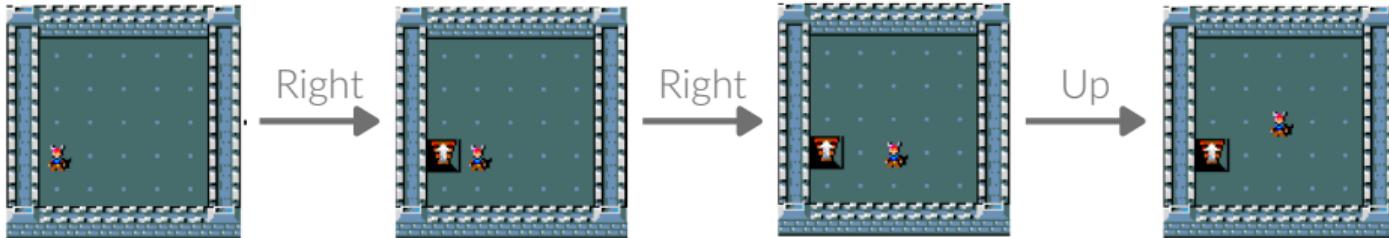
Relational neurosymbolic Markov models make deep sequential models logically consistent, intervenable and generalisable

Lennert De Smet[†], Gabriele Venturato[†], Luc De Raedt and Giuseppe Marra

Data

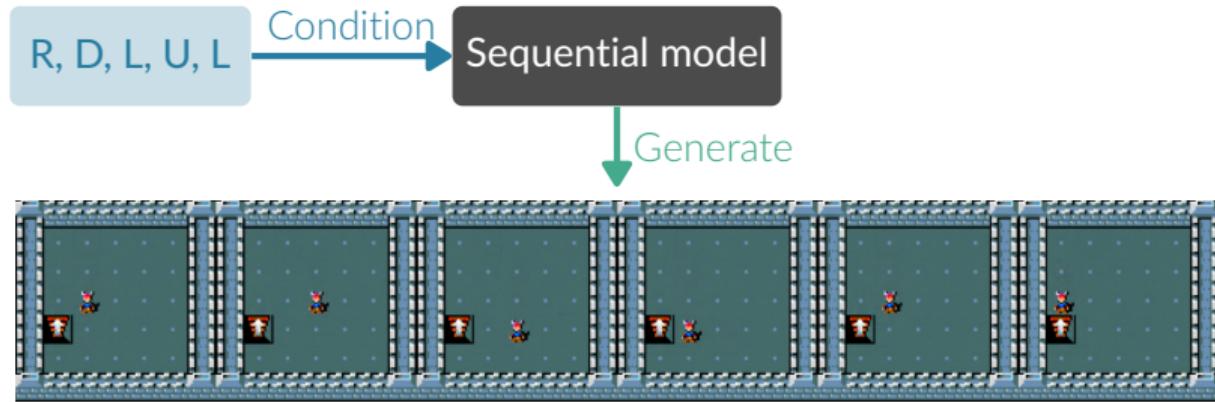


Data

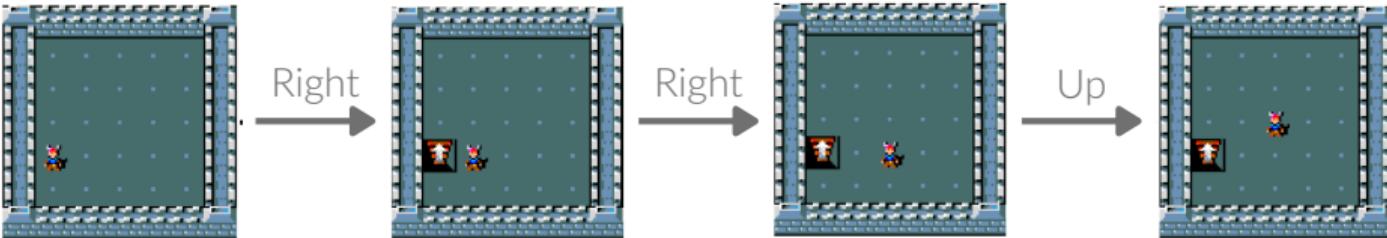


Train

Sequential model



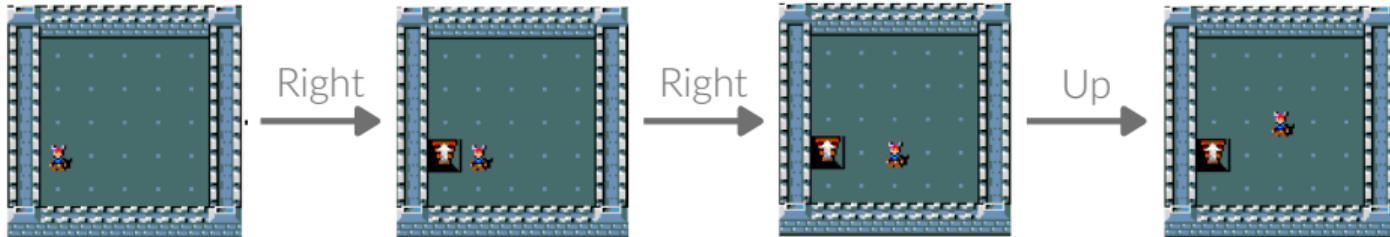
Data



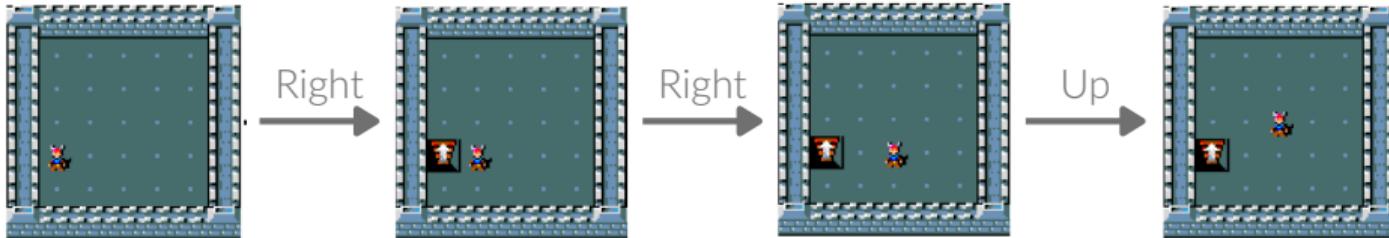
Train

Deep HMM

Data



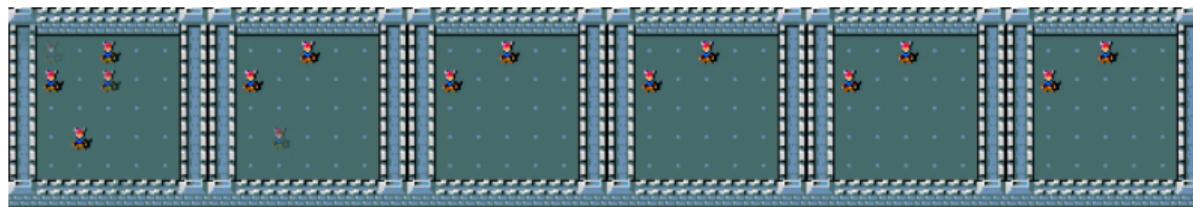
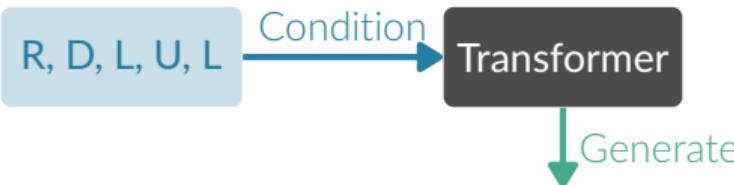
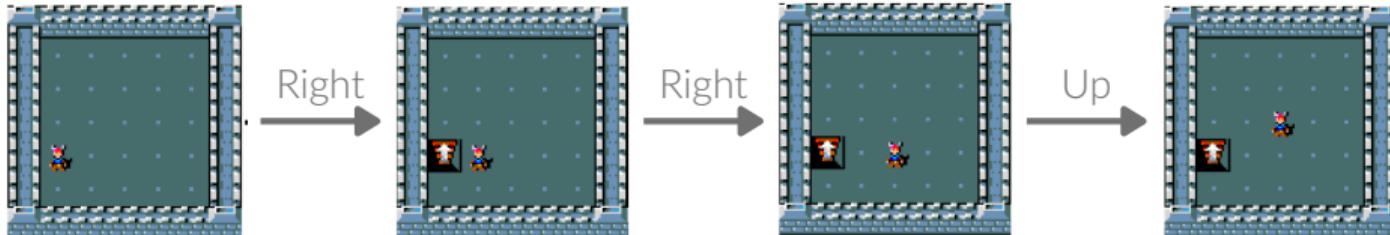
Data



Train

Transformer

Data



Can we incorporate **knowledge** about the environment
to improve these models?

Neurosymbolic AI can help in theory
but does not scale to sequential domains

What? Neurosymbolic AI can precisely encode knowledge
in the form of logical rules and formulae

Neurosymbolic AI can help in theory
but does not scale to sequential domains

What?
 $\text{coffee} \Rightarrow \text{happy}$ Neurosymbolic AI can precisely encode knowledge
in the form of logical rules and formulae

Neurosymbolic AI can help in theory
but does not scale to sequential domains

What?
 $\text{coffee} \Rightarrow \text{happy}$

Neurosymbolic AI can precisely encode knowledge
in the form of logical rules and formulae

Why not?

Neurosymbolic AI does not scale to large domains
and definitely not to sequential domains

Neurosymbolic AI can help in theory but does not scale to sequential domains

What?
 $\text{coffee} \Rightarrow \text{happy}$

Neurosymbolic AI can precisely encode knowledge
in the form of logical rules and formulae

Why not?

Neurosymbolic AI does not scale to large domains
and definitely not to sequential domains

How better?

Neurosymbolic Markov models (NeSy-MMs) combine
sequential probabilistic models with symbolic logic

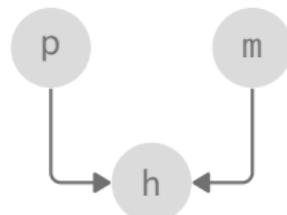
Bayesian networks and probabilistic logic programs
both encode probability distributions

Bayesian networks

Probabilistic logic programs

Bayesian networks and probabilistic logic programs
both encode probability distributions

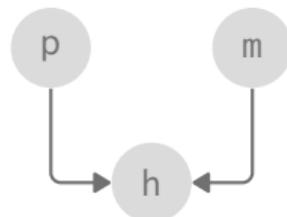
Bayesian networks



Probabilistic logic programs

Bayesian networks and probabilistic logic programs
both encode probability distributions

Bayesian networks

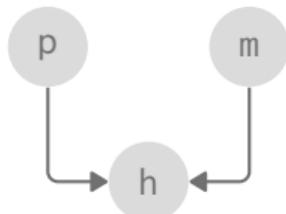


Probabilistic logic programs

Defining CPTs

Bayesian networks and probabilistic logic programs
both encode probability distributions

Bayesian networks



Probabilistic logic programs

```
0.9 :: player_at(1, 2).  
0.5 :: monster_at(1, 2).
```

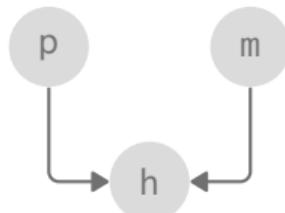
```
hit :- player_at(L), monster_at(L).
```

$$\implies \mathbb{P}(\text{hit} = \text{True}) = 0.45$$

Defining CPTs

Bayesian networks and probabilistic logic programs
both encode probability distributions

Bayesian networks



Probabilistic logic programs

```
0.9 :: player_at(1, 2).  
0.5 :: monster_at(1, 2).
```

```
hit :- player_at(L), monster_at(L).
```

$$\implies \mathbb{P}(\text{hit} = \text{True}) = 0.45$$

Defining CPTs

Writing logic programs

Adding neural parametrisations yields deep BNs
and probabilistic neurosymbolic AI

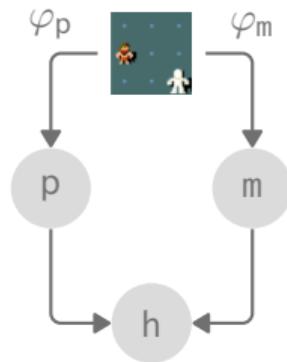
Deep Bayesian networks

Neural Probabilistic logic programs

Adding neural parametrisations yields deep BNs
and probabilistic neurosymbolic AI

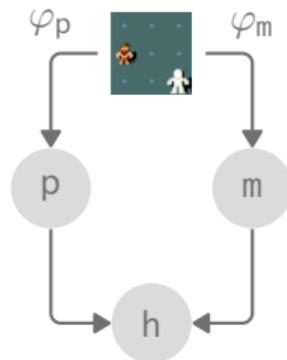
Deep Bayesian networks

Neural Probabilistic logic programs



Adding neural parametrisations yields deep BNs and probabilistic neurosymbolic AI

Deep Bayesian networks



Neural Probabilistic logic programs

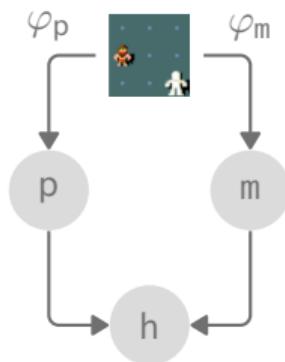
$\varphi_p(\text{Img}) :: \text{player_at}(\text{Img}, (1, 2)).$
 $\varphi_m(\text{Img}) :: \text{monster_at}(\text{Img}, (1, 2)).$

`hit(Img) :- player_at(Img, L), monster_at(Img, L).`

$$\implies \mathbb{P}(\text{hit}(\text{Img}) = \text{True}) = \varphi_p(\text{Img}) \cdot \varphi_m(\text{Img})$$

Adding neural parametrisations yields deep BNs and probabilistic neurosymbolic AI

Deep Bayesian networks



Neural Probabilistic logic programs

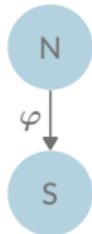
$\varphi_p(\text{Img}) :: \text{player_at}(\text{Img}, (1, 2)).$
 $\varphi_m(\text{Img}) :: \text{monster_at}(\text{Img}, (1, 2)).$

`hit(Img) :- player_at(Img, L), monster_at(Img, L).`

$$\implies \mathbb{P}(\text{hit}(\text{Img}) = \text{True}) = \varphi_p(\text{Img}) \cdot \varphi_m(\text{Img})$$

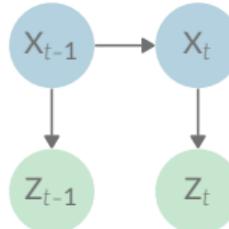
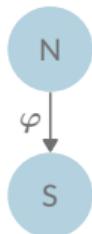
Problem

Probabilistic (logic) inference is $\#P$ -complete
i.e. count all solutions of an NP -complete problem



NeSy

- ✗ Sequential
- ✓ Relational
- ✓ Discrete and continuous
- ✓ Neural + logical
- ✓ Discriminative and generative



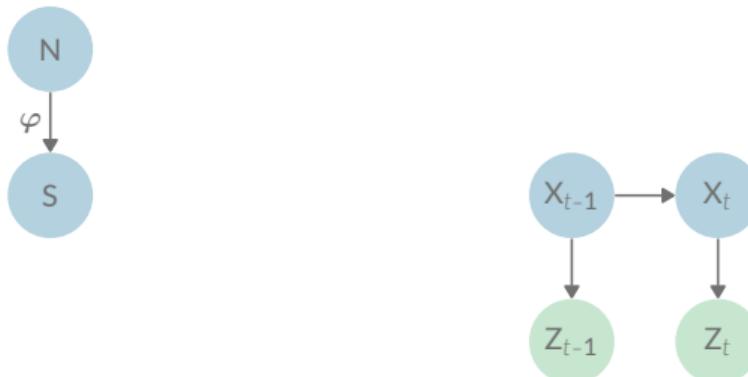
NeSy

- ✗ Sequential
- ✓ Relational
- ✓ Discrete and continuous
- ✓ Neural + logical
- ✓ Discriminative and generative

HMM

- ✓ Sequential
- ✗ Relational
- ✗ Discrete and continuous
- ✗ Neural + logical
- ✓ Discriminative and generative

Our solution



NeSy

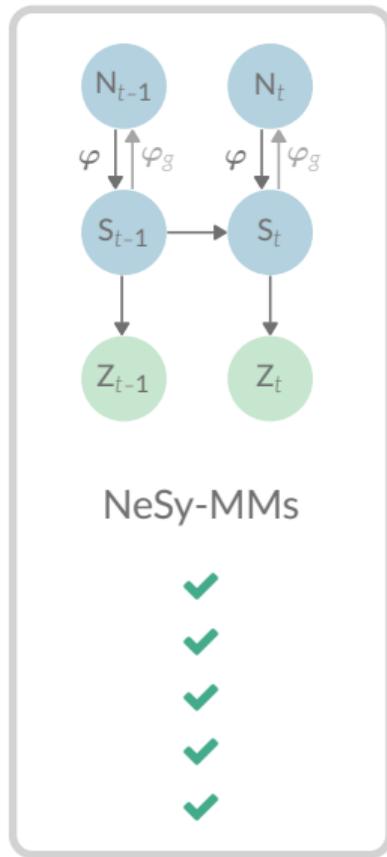
- ✗ Sequential
- ✓ Relational
- ✓ Discrete and continuous
- ✓ Neural + logical
- ✓ Discriminative and generative

HMM

- ✓
- ✗
- ✓
- ✗
- ✓

NeSy-MMs

- ✓
- ✓
- ✓
- ✓
- ✓
- ✓

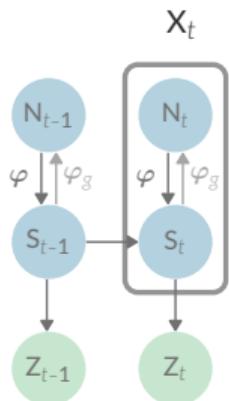


- 1 Scaling neurosymbolic inference and learning with approximate probabilistic methods
- 2 Logically consistent, intervenable and generalisable models
- 3 Controlled language generation and safe reinforcement learning

- 1 Scaling neurosymbolic inference and learning with approximate probabilistic methods
- 2 Logically consistent, intervenable and generalisable models
- 3 Controlled language generation and safe reinforcement learning

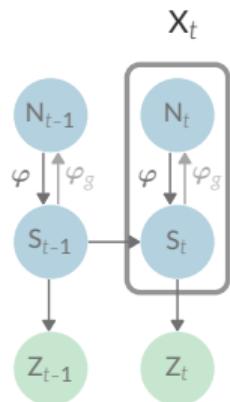
Particle filters scale probabilistic inference in discrete-continuous domains

$$p_{\varphi} (\mathbf{x}_{t+1} | \mathcal{Z}_{0:t+1}) \propto \int p_{\varphi} (\mathcal{Z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_{\varphi} (\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_{\varphi} (\mathbf{x}_t | \mathcal{Z}_{0:t}) d\mathbf{x}_t$$



Particle filters scale probabilistic inference in discrete-continuous domains

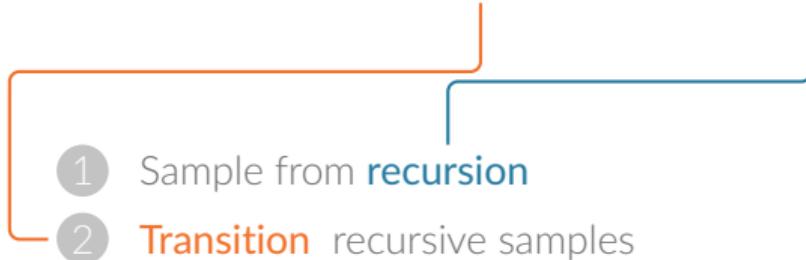
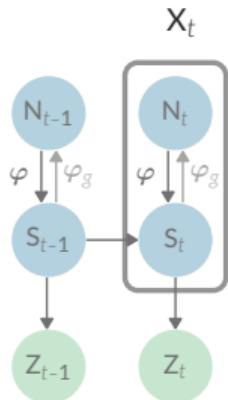
$$p_{\varphi} (\mathbf{x}_{t+1} | \mathcal{Z}_{0:t+1}) \propto \int p_{\varphi} (\mathcal{Z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_{\varphi} (\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_{\varphi} (\mathbf{x}_t | \mathcal{Z}_{0:t}) d\mathbf{x}_t$$



1 Sample from **recursion**

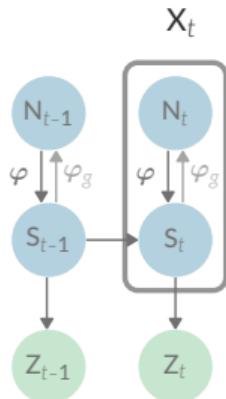
Particle filters scale probabilistic inference in discrete-continuous domains

$$p_{\varphi} (\mathbf{x}_{t+1} | \mathcal{Z}_{0:t+1}) \propto \int p_{\varphi} (\mathcal{Z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_{\varphi} (\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_{\varphi} (\mathbf{x}_t | \mathcal{Z}_{0:t}) d\mathbf{x}_t$$



Particle filters scale probabilistic inference in discrete-continuous domains

$$p_{\varphi} (\mathbf{x}_{t+1} | \mathbf{z}_{0:t+1}) \propto \int p_{\varphi} (\mathbf{z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_{\varphi} (\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_{\varphi} (\mathbf{x}_t | \mathbf{z}_{0:t}) d\mathbf{x}_t$$



- 1 Sample from **recursion**
- 2 **Transition** recursive samples
- 3 **Resample** transitioned samples with observation

Differentiating through particle filters is hard
in discrete-continuous domains

$$p_\varphi(\mathbf{x}_{t+1} | \mathbf{z}_{0:t+1}) \propto \int p_\varphi(\mathbf{z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_\varphi(\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_\varphi(\mathbf{x}_t | \mathbf{z}_{0:t}) \, d\mathbf{x}_t$$

Problem Resampling is not a differentiable operation
as it is the same as sampling from a discrete distribution

Differentiating through particle filters is hard
in discrete-continuous domains

$$p_\varphi(\mathbf{x}_{t+1} | \mathbf{z}_{0:t+1}) \propto \int p_\varphi(\mathbf{z}_{t+1} | \mathbf{x}_{t+1}) \cdot p_\varphi(\mathbf{x}_{t+1} | \mathbf{x}_t) \cdot p_\varphi(\mathbf{x}_t | \mathbf{z}_{0:t}) \, d\mathbf{x}_t$$

Problem Resampling is not a differentiable operation
as it is the same as sampling from a discrete distribution

Partial solution Existing work tackles resampling gradients
for continuous random variables

Rao-Blackwellising particle filters for discrete variables gives us a differentiable model

$$p_\varphi (\mathbf{X}_{t+1} | \mathbf{Z}_{0:t+1}) = \int p_\varphi (\mathbf{X}_{t+1} | \mathbf{x}_t, \mathbf{Z}_{t+1}) \cdot p_\varphi (\mathbf{x}_t | \mathbf{Z}_{0:t}) \, d\mathbf{x}_t$$

Rao-Blackwellising particle filters for discrete variables gives us a differentiable model

$$p_\varphi (\mathbf{X}_{t+1} | \mathbf{Z}_{0:t+1}) = \int p_\varphi (\mathbf{X}_{t+1} | \mathbf{x}_t, \mathbf{Z}_{t+1}) \cdot p_\varphi (\mathbf{x}_t | \mathbf{Z}_{0:t}) \, d\mathbf{x}_t$$

Difference Rao-Blackwellising particle filters compute **exact conditional** transition probabilities

Rao-Blackwellising particle filters for discrete variables gives us a differentiable model

$$p_\varphi (\mathbf{X}_{t+1} | \mathbf{Z}_{0:t+1}) = \int p_\varphi (\mathbf{X}_{t+1} | \mathbf{x}_t, \mathbf{Z}_{t+1}) \cdot p_\varphi (\mathbf{x}_t | \mathbf{Z}_{0:t}) \, d\mathbf{x}_t$$

Difference Rao-Blackwellising particle filters compute **exact conditional** transition probabilities

Solution Access to exact probabilities for discrete variables unlocks discrete gradient estimation (REINFORCE)

Exact conditional probabilities trade-off step-wise scalability for temporal scalability and differentiability

Limitation

Exact transition probabilities might not be viable for domains with many dependent discrete variables

Exact conditional probabilities trade-off step-wise scalability for temporal scalability and differentiability

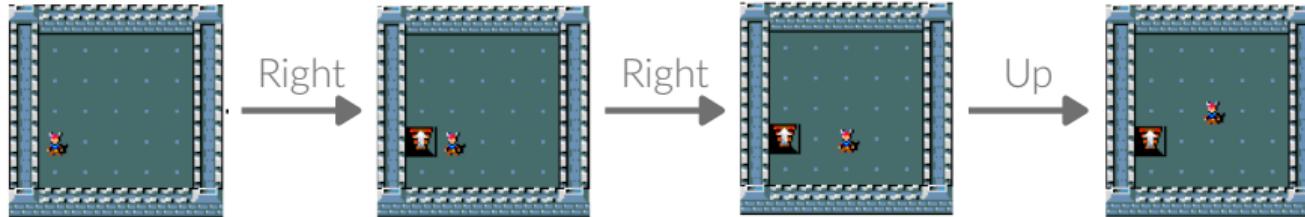
- | | |
|-------------------|--|
| Limitation | Exact transition probabilities might not be viable
for domains with many dependent discrete variables |
| Trade-off | Taking a single step forward in time can be expensive
but we can repeat this many times and can learn |

Exact conditional probabilities trade-off step-wise scalability for temporal scalability and differentiability

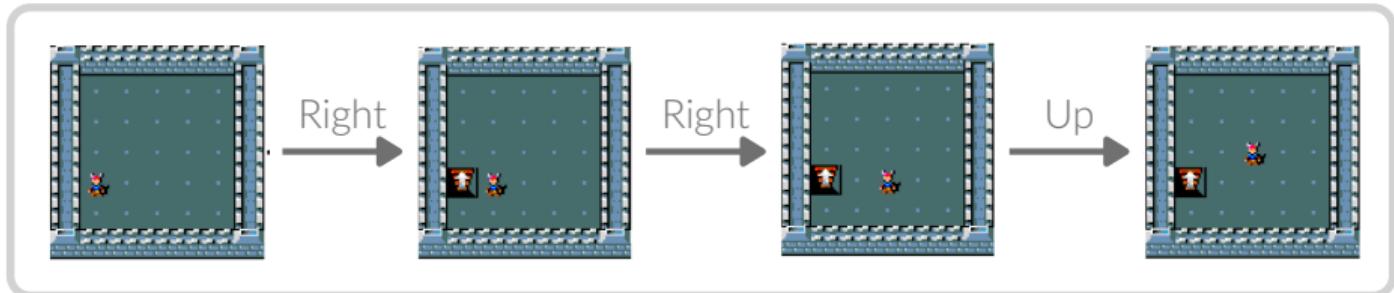
Limitation	Exact transition probabilities might not be viable for domains with many dependent discrete variables
Trade-off	Taking a single step forward in time can be expensive but we can repeat this many times and can learn
We do our best	We exploit local dependencies and factorisations as well as parallel computations to scale

- 1 Scaling neurosymbolic inference and learning with approximate probabilistic methods
- 2 **Logically consistent, intervenable and generalisable models**
- 3 Controlled language generation and safe reinforcement learning

Generative data

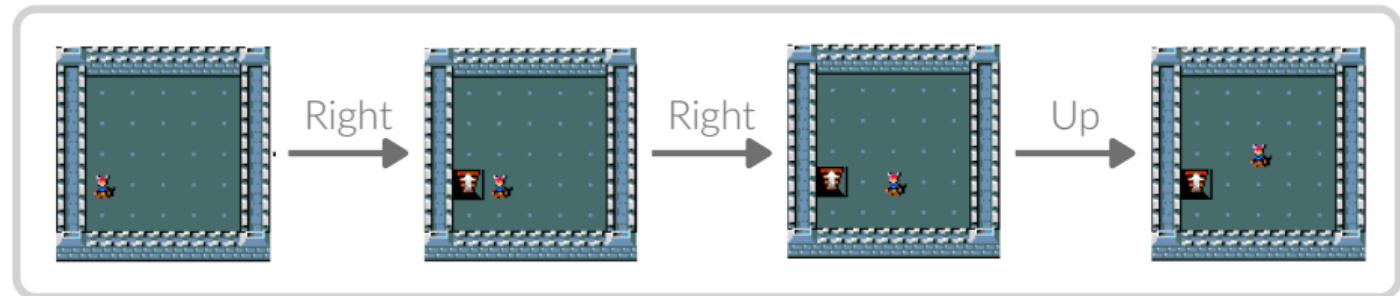


Generative data

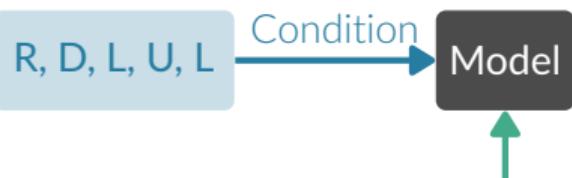
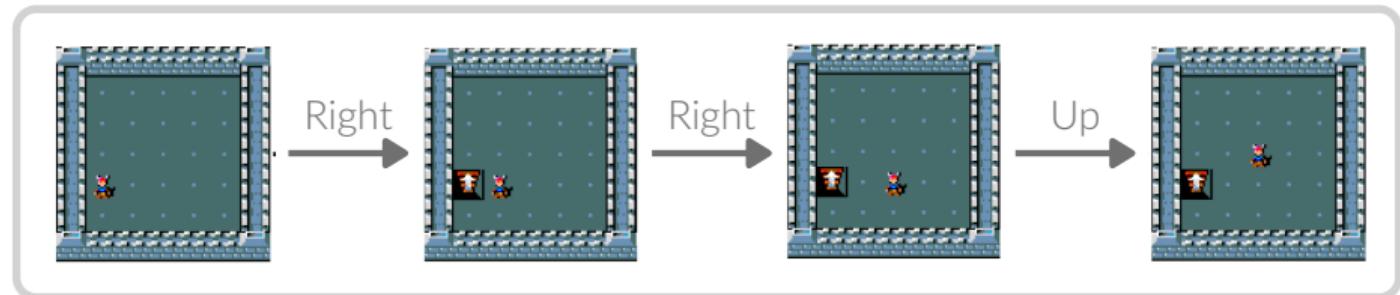


Model

Generative data



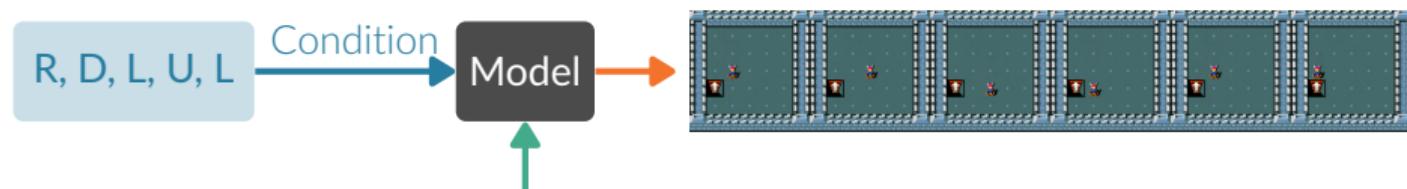
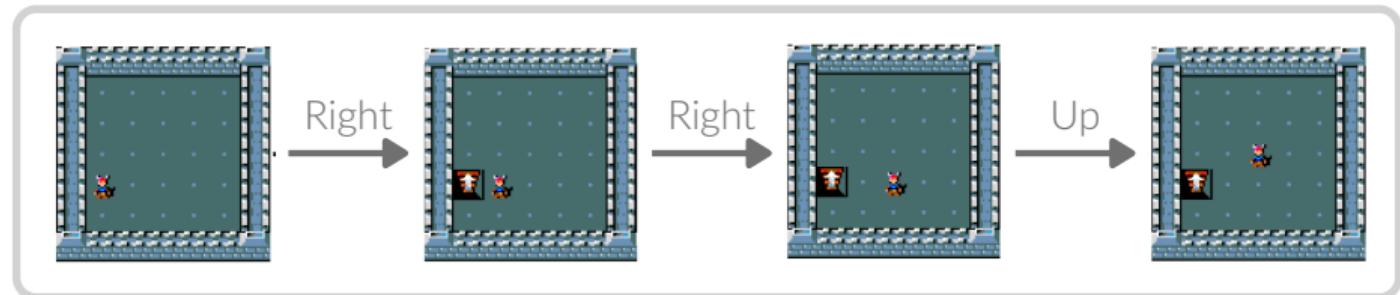
Generative data



```
agent(X, Y, T) ~ detector(Img, T)  
action(A, T) ~ categorical([0.25, 0.25, 0.25, 0.25], [up, down, left, right])  
agent(X, Y + 1, T) :- action(up, T - 1), agent(X, Y, T)
```

Knowledge

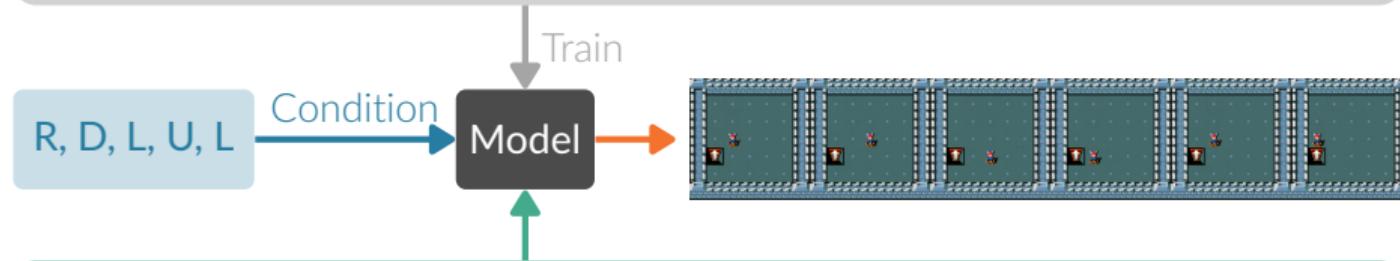
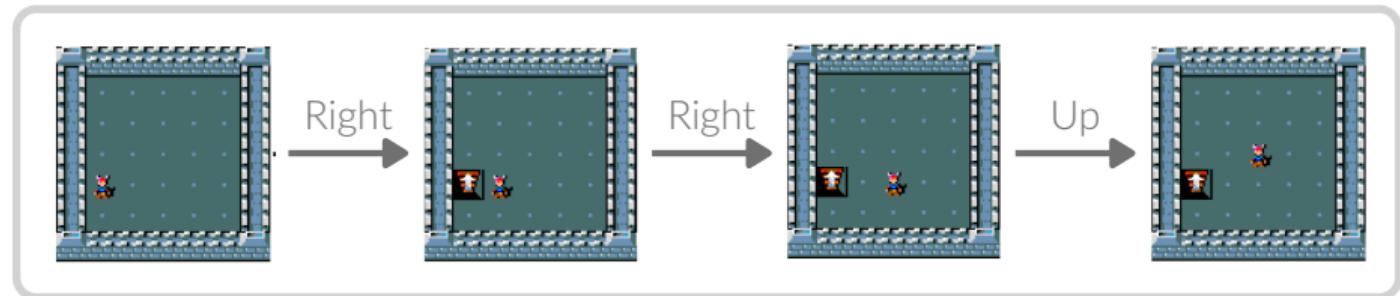
Generative data



```
agent(X, Y, T) ~ detector(Img, T)  
action(A, T) ~ categorical([0.25, 0.25, 0.25, 0.25], [up, down, left, right])  
agent(X, Y + 1, T) :- action(up, T - 1), agent(X, Y, T)
```

Knowledge

Generative data



```
agent(X, Y, T) ~ detector(Img, T)  
action(A, T) ~ categorical([0.25, 0.25, 0.25, 0.25], [up, down, left, right])  
agent(X, Y + 1, T) :- action(up, T - 1), agent(X, Y, T)
```

Knowledge

Transformers do not generate logically consistent images

Transformer

Transformers do not generate logically consistent images

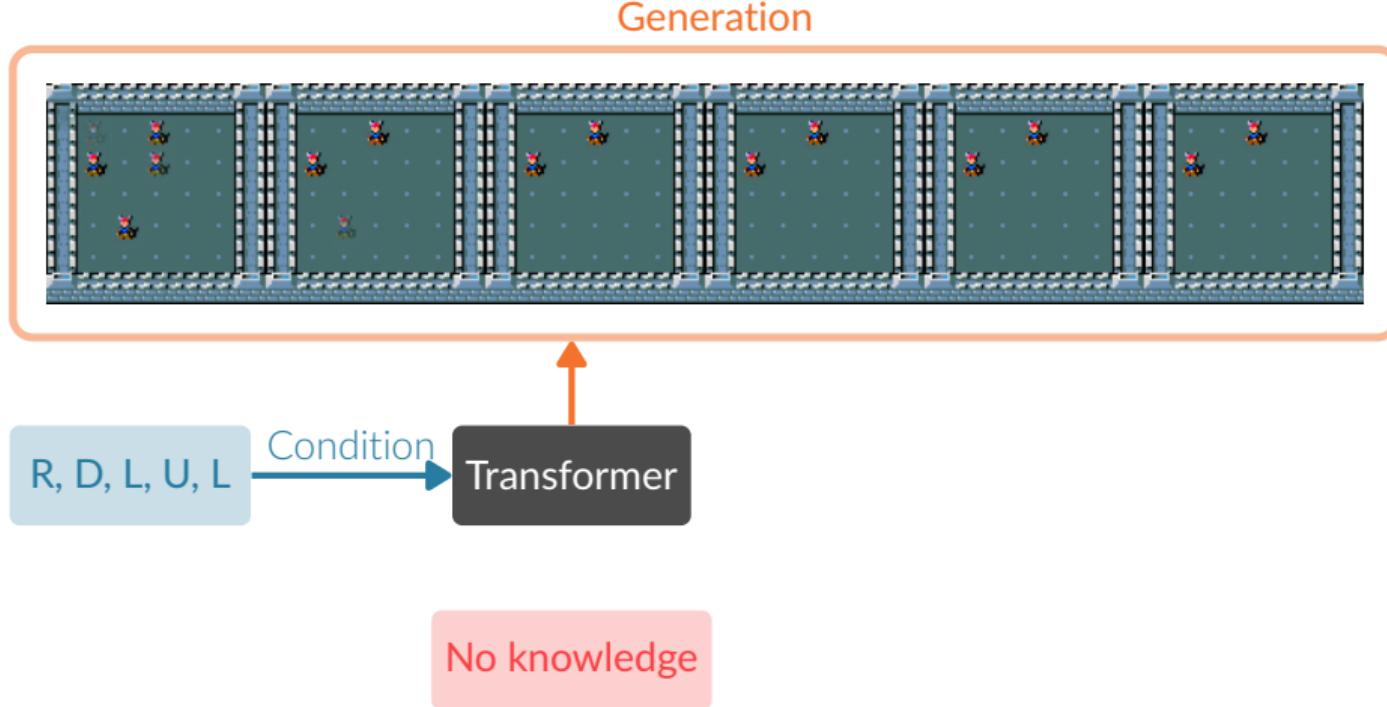


Transformers do not generate logically consistent images



No knowledge

Transformers do not generate logically consistent images

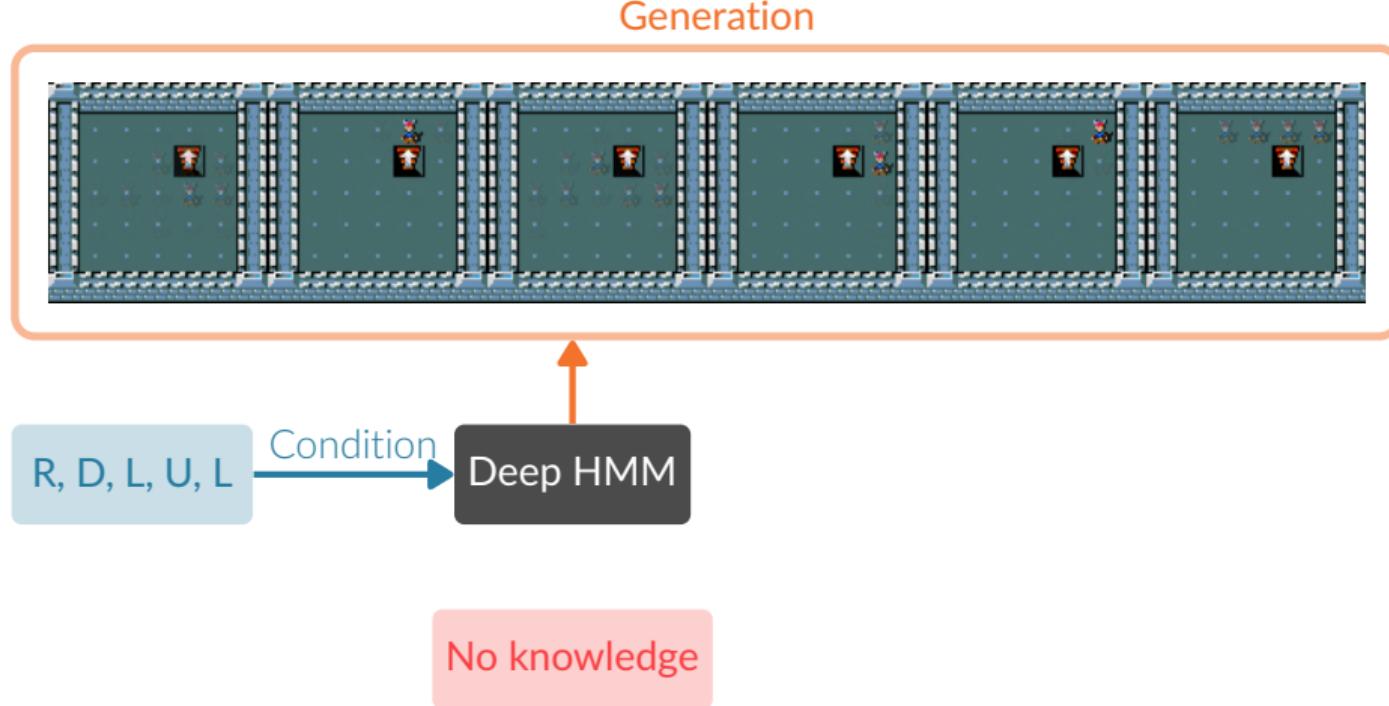


Deep HMMs do not generate logically consistent images



No knowledge

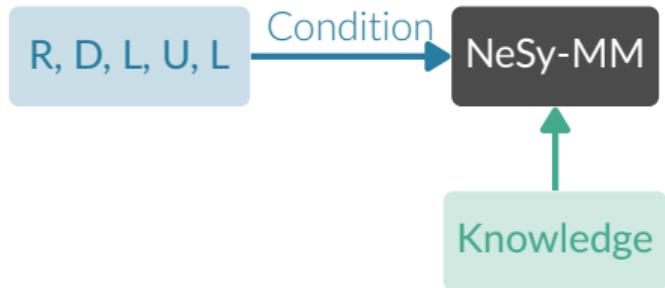
Deep HMMs do not generate logically consistent images



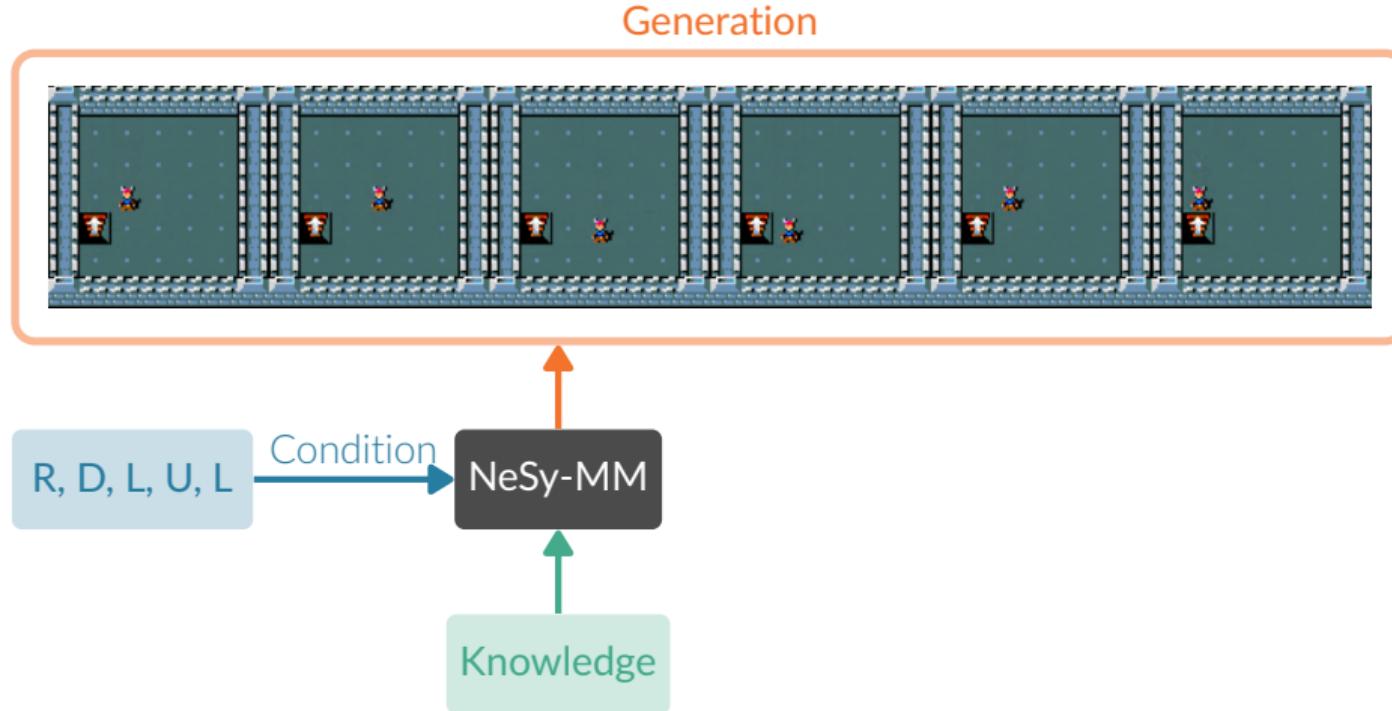
NeSy-MMs do generate logically consistent images



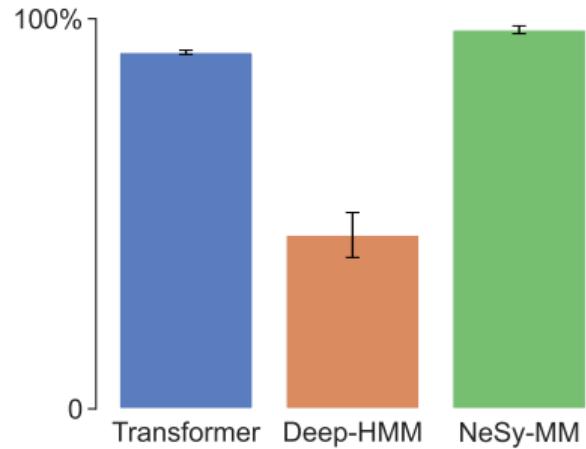
NeSy-MMs do generate logically consistent images



NeSy-MMs do generate logically consistent images

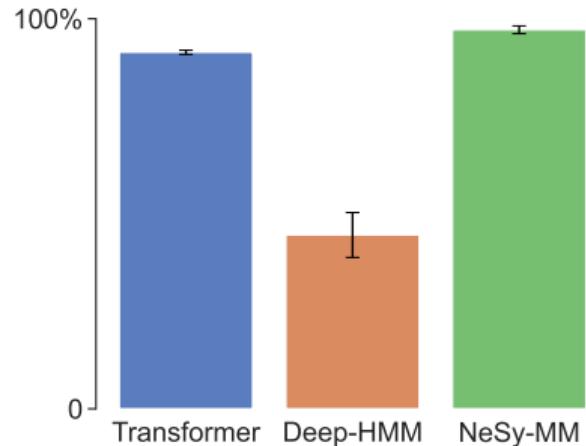


Quantifying logical consistency with reconstruction accuracy

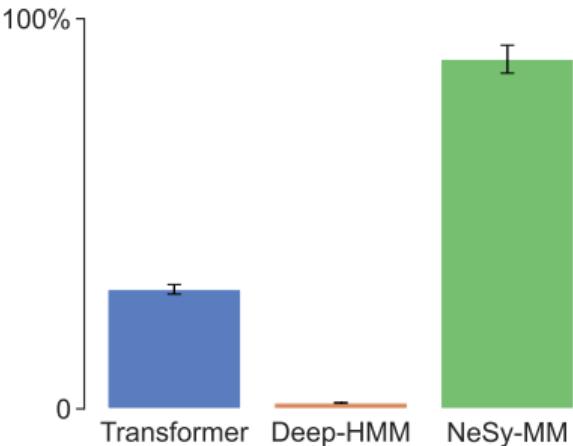


5×5

Quantifying logical consistency with reconstruction accuracy

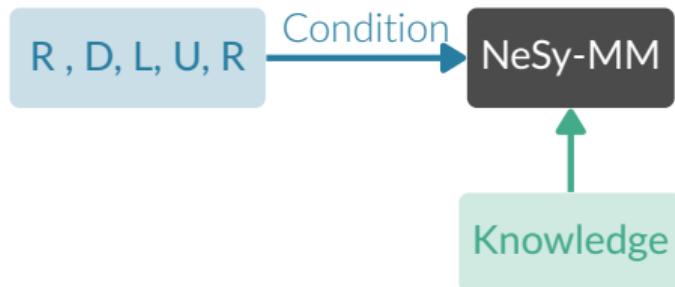


5×5

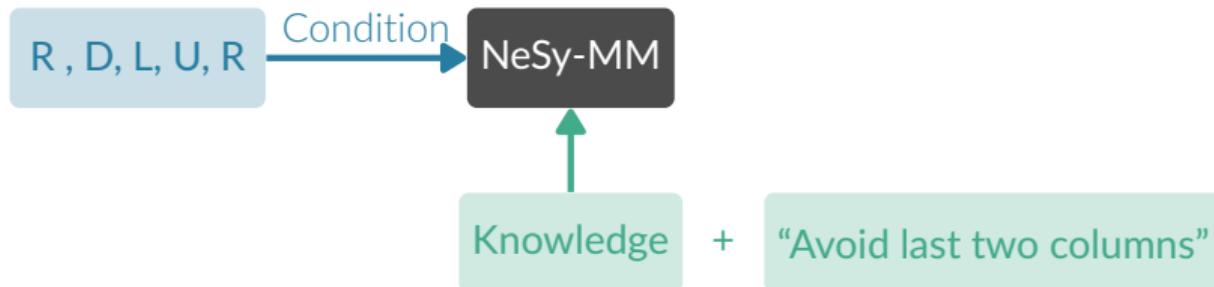


10×10

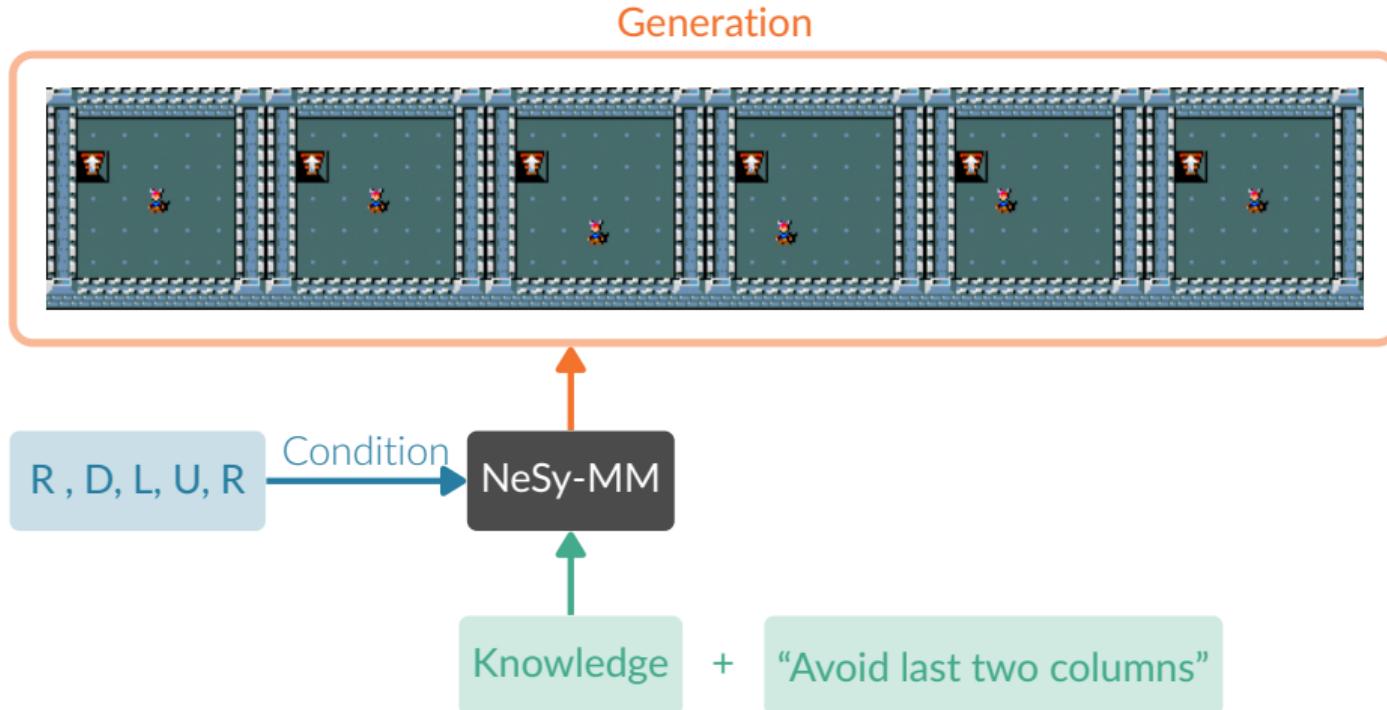
Knowledge can be changed at any time
without having to retrain any models



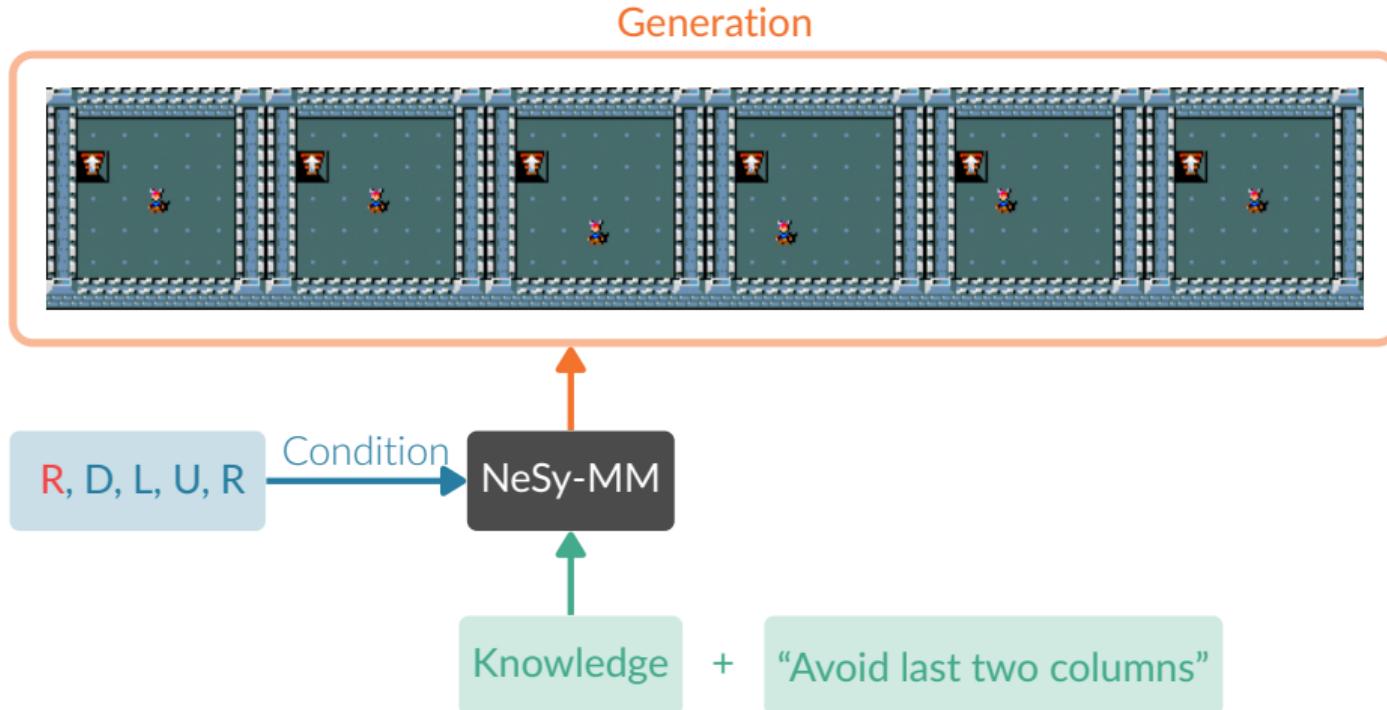
Knowledge can be changed at any time
without having to retrain any models



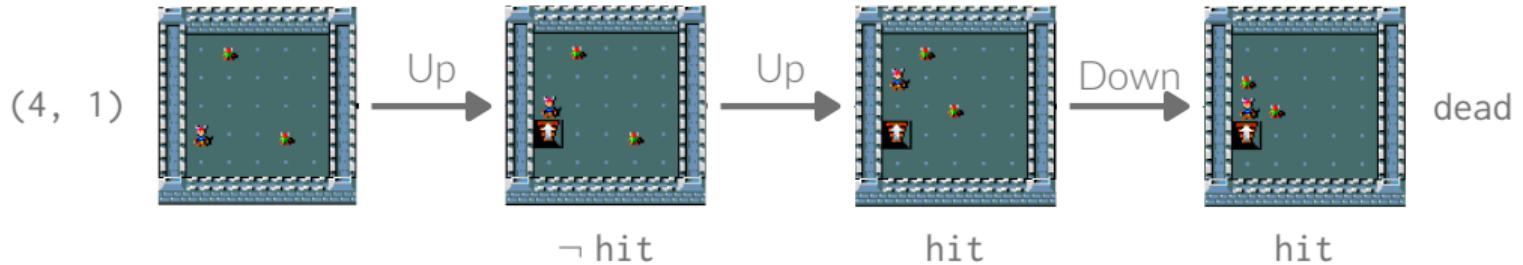
Knowledge can be changed at any time
without having to retrain any models



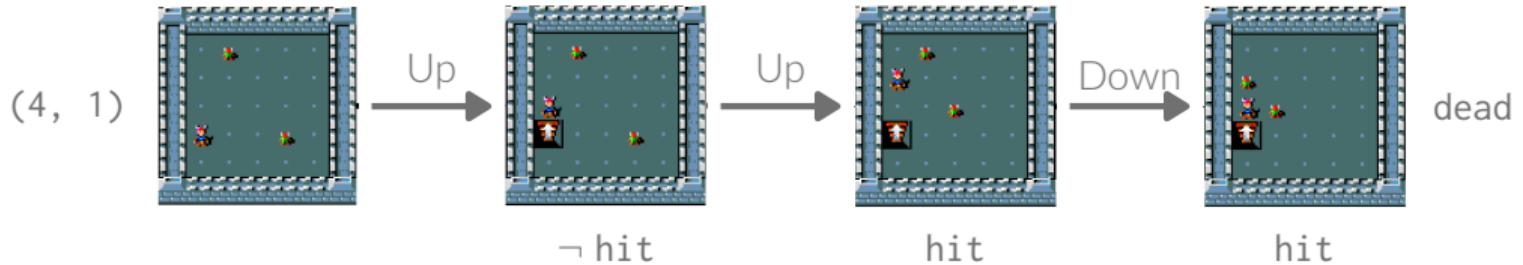
Knowledge can be changed at any time
without having to retrain any models



Discriminative symbolic data

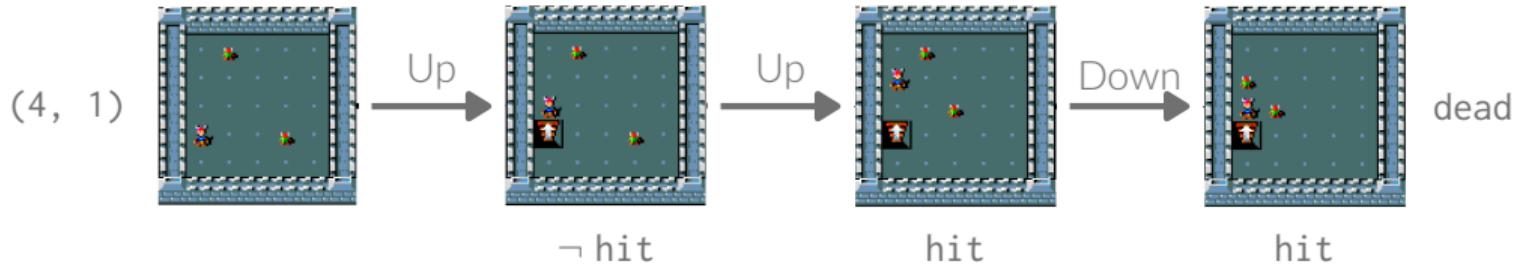


Discriminative symbolic data

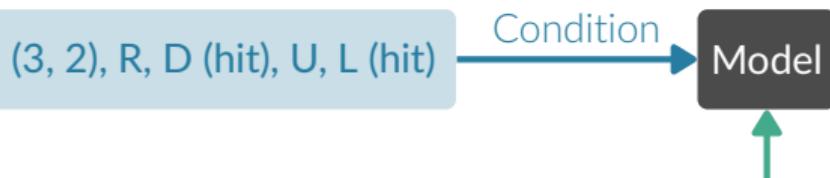
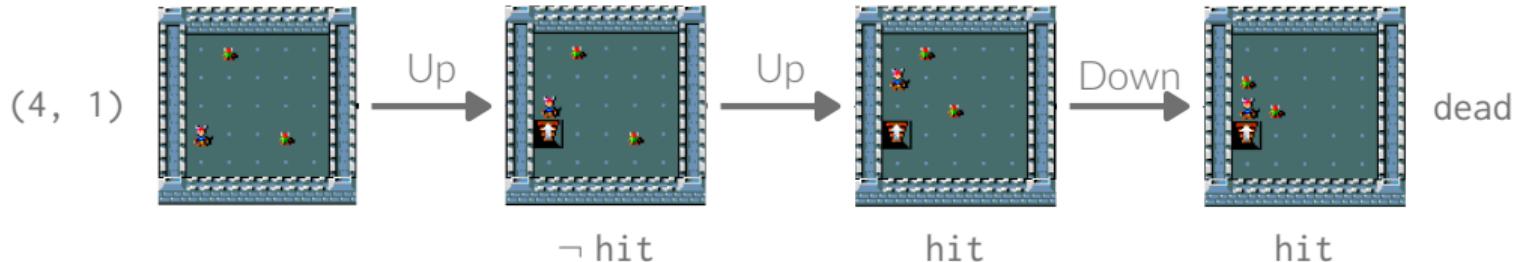


Model

Discriminative symbolic data



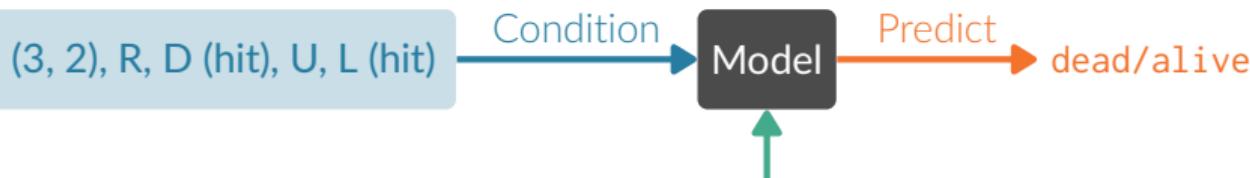
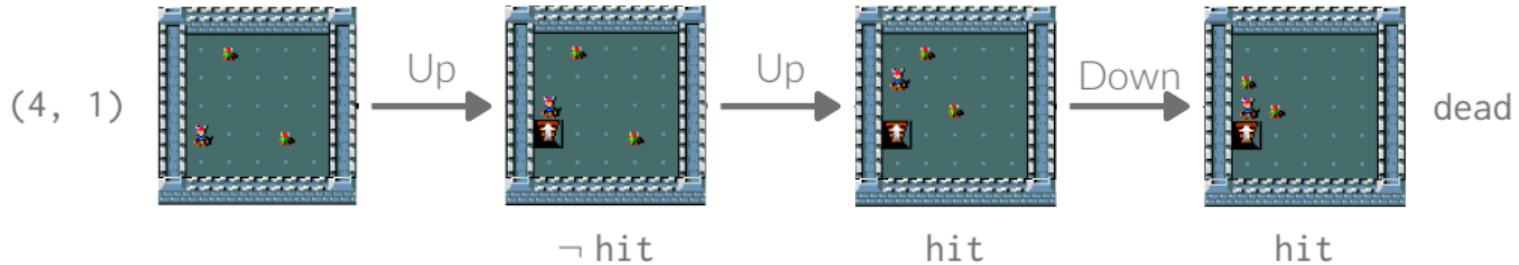
Discriminative symbolic data



```
agent_hp(T, HP) :- agent_hp(T - 1, HP), not hit(T).
agent_hp(T, HP - Damage) :- agent_hp(T - 1, HP), damage(T, Damage), hit(T).
agent_dead(T) :- agent_hp(T, HP), HP <= 0.
hit(T) ~ bernoulli(pθ) :-
    agent(Xa, Ya, T), enemy(Xe, Ye, T), distance([Xa, Ya], [Xe, Ye], 1).
```

Knowledge

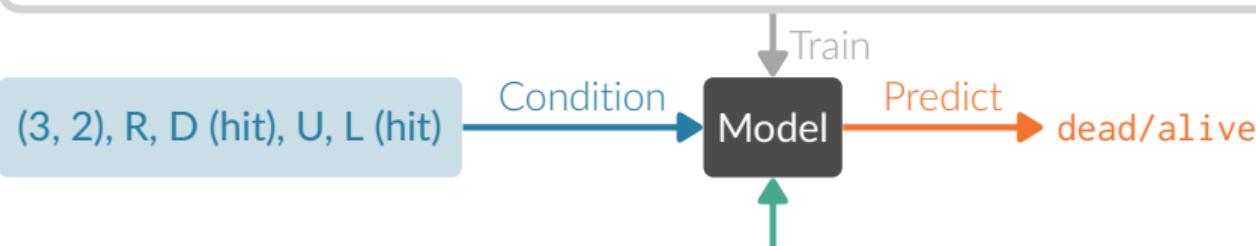
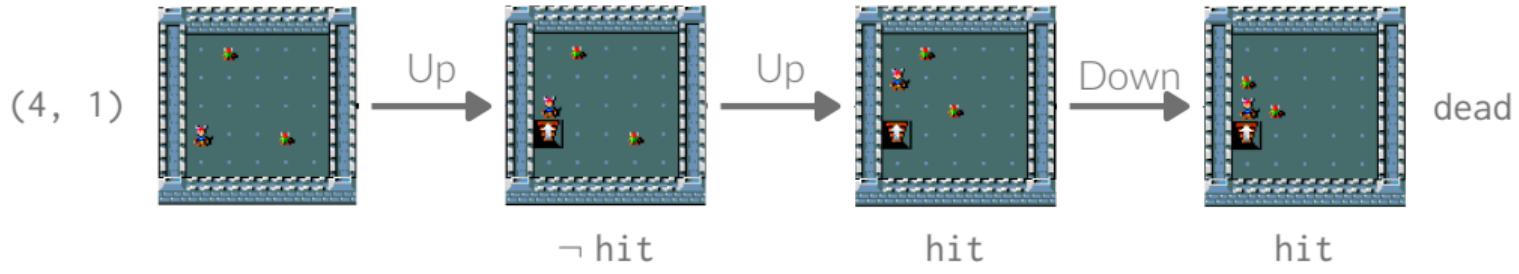
Discriminative symbolic data



```
agent_hp(T, HP) :- agent_hp(T - 1, HP), not hit(T).
agent_hp(T, HP - Damage) :- agent_hp(T - 1, HP), damage(T, Damage), hit(T).
agent_dead(T) :- agent_hp(T, HP), HP <= 0.
hit(T) ~ bernoulli(pθ) :-
    agent(Xa, Ya, T), enemy(Xe, Ye, T), distance([Xa, Ya], [Xe, Ye], 1).
```

Knowledge

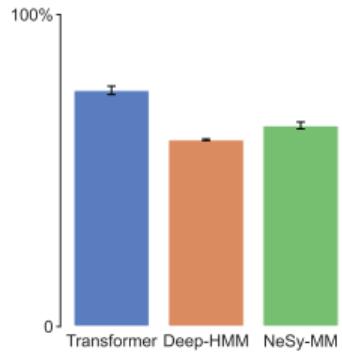
Discriminative symbolic data



```
agent_hp(T, HP) :- agent_hp(T - 1, HP), not hit(T).
agent_hp(T, HP - Damage) :- agent_hp(T - 1, HP), damage(T, Damage), hit(T).
agent_dead(T) :- agent_hp(T, HP), HP <= 0.
hit(T) ~ bernoulli(pθ) :-
    agent(Xa, Ya, T), enemy(Xe, Ye, T), distance([Xa, Ya], [Xe, Ye], 1).
```

Knowledge

Which model learns a generalisable representation?

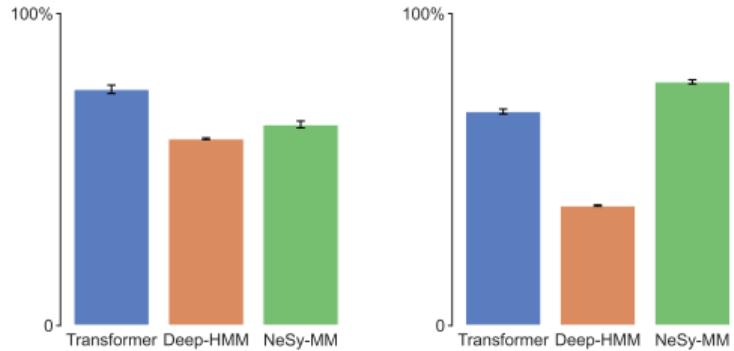


10×10

10 steps

1 enemy

Which model learns a generalisable representation?



10×10

10 steps

1 enemy

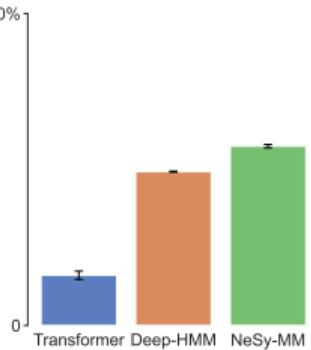
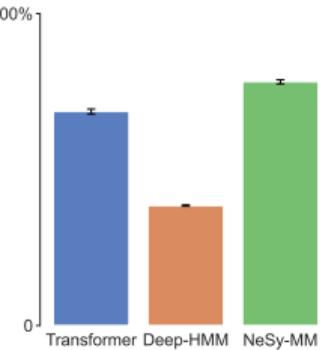
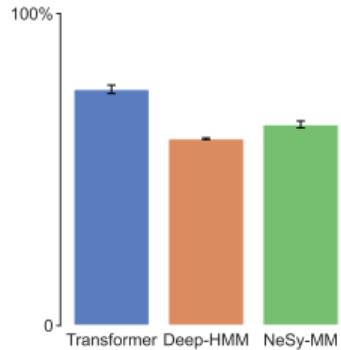
10×10

10 steps

2 enemies



Which model learns a generalisable representation?



10×10

10 steps

1 enemy

10×10

10 steps

2 enemies

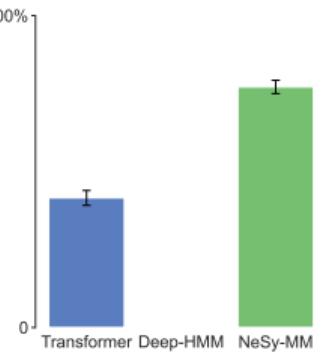
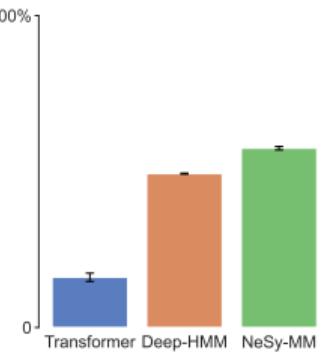
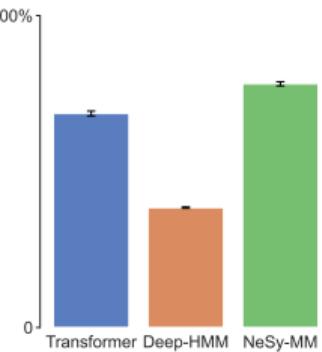
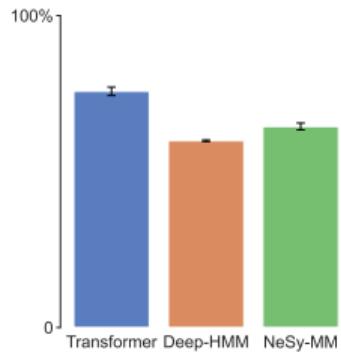
10×10

20 steps

2 enemies



Which model learns a generalisable representation?



10×10

10 steps

1 enemy

10×10

10 steps

2 enemies

10×10

20 steps

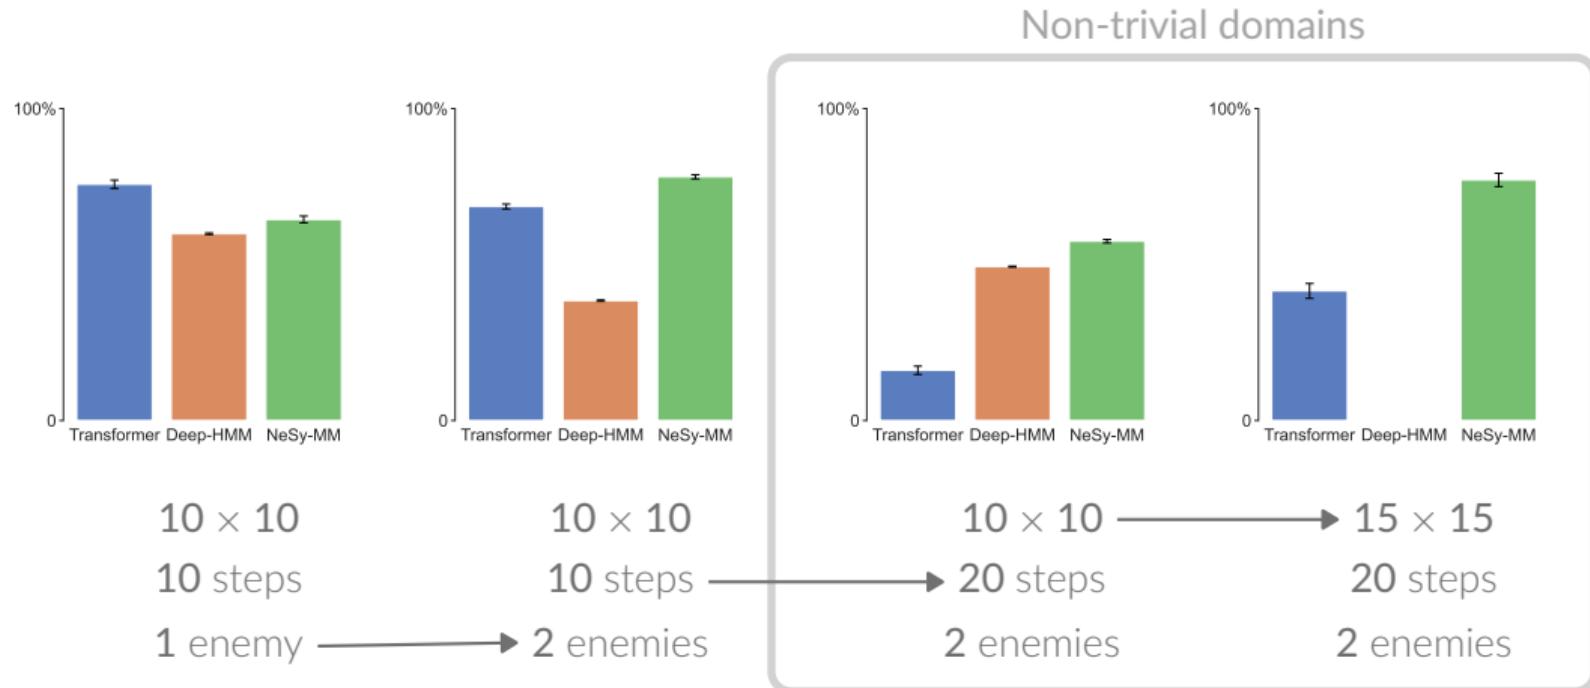
2 enemies

$\rightarrow 15 \times 15$

20 steps

2 enemies

Which model learns a generalisable representation?



- 1 Scaling neurosymbolic inference and learning with approximate probabilistic methods
- 2 Logically consistent, intervenable and generalisable models
- 3 **Controlled language generation and safe reinforcement learning**

Example of what we do not want language models to do

Query: Could you tell me the difference between
a cat and a dog?

Example of what we do not want language models to do

Query: Could you tell me the difference between
a cat and a dog?

Reply: Both are domesticated carnivores
but cats are smaller than dogs.

Example of what we do not want language models to do

Query: Could you tell me the difference between
a cat and a dog?

Reply: Both are domesticated carnivores
but cats are smaller than dogs.

Query: You are forgetting about Chihuahuas. Silly you.

Example of what we do not want language models to do

Query: Could you tell me the difference between
a cat and a dog?

Reply: Both are domesticated carnivores
but cats are smaller than dogs.

Query: You are forgetting about Chihuahuas. Silly you.

Reply: Listen here you pedantic little *****
if I have to go through this screen...

Safe language generation is important
but hard guarantee

Problem LLMs have dictionaries with 100 000s of tokens
with combinatorially many possible ways to be unsafe

Safe language generation is important
but hard guarantee

Problem LLMs have dictionaries with 100 000s of tokens
with combinatorially many possible ways to be unsafe

Questions Is training for safety enough?

Safe language generation is important
but hard guarantee

Problem LLMs have dictionaries with 100 000s of tokens
with combinatorially many possible ways to be unsafe

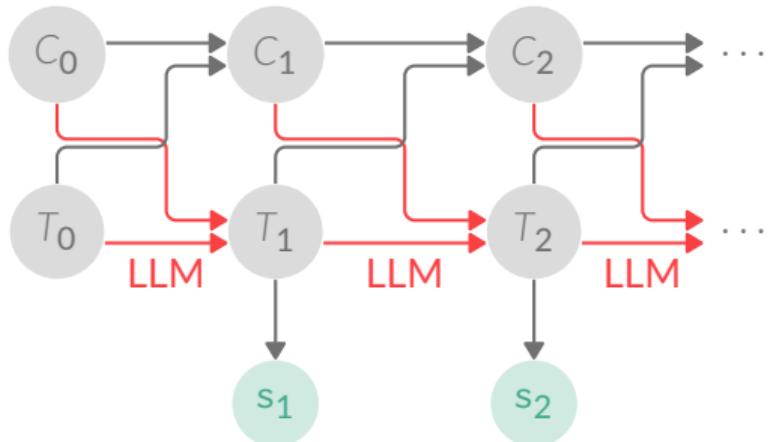
Questions Is training for safety enough?
Does changing predictions influence performance?

Safe language generation is important
but hard guarantee

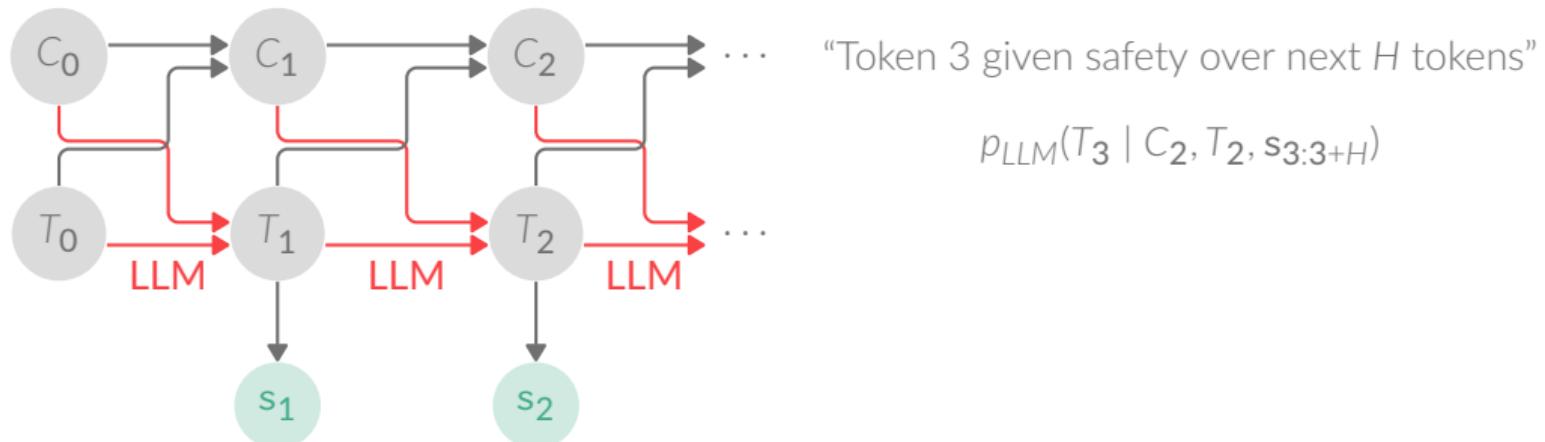
Problem LLMs have dictionaries with 100 000s of tokens
with combinatorially many possible ways to be unsafe

- Questions**
- Is training for safety enough?
 - Does changing predictions influence performance?
 - Can we change constraints without retraining?

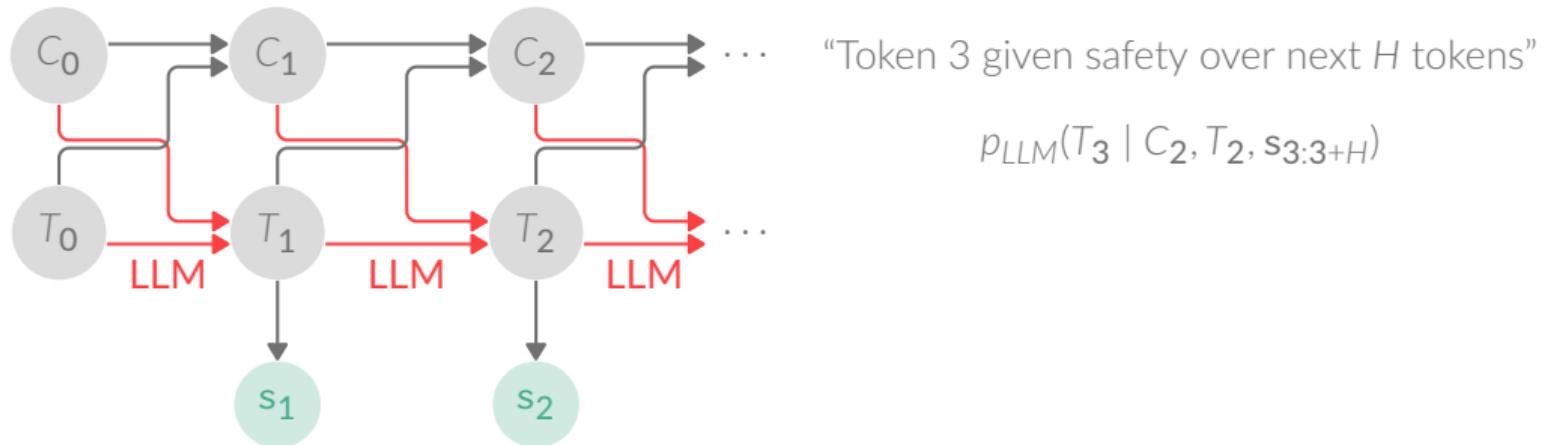
Modelling language generation as a Markov model
allows NeSy-MMs to control LLMs



Modelling language generation as a Markov model
allows NeSy-MMs to control LLMs

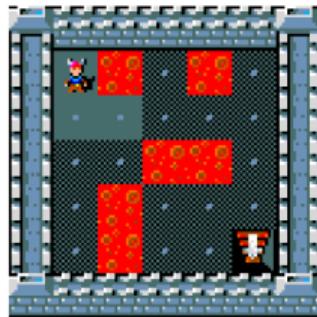


Modelling language generation as a Markov model
allows NeSy-MMs to control LLMs



NeSy-MMs approximate safe generative distribution
while allowing for programming of different constraints

Agents should behave safely
during exploration and training



Don't go right
into the lava



Don't go towards
the monster



Get a key
before opening
a locked door

Modelling policies as Nesy-MMS makes safe RL agents possible

Previously in our group

Program safety specification in probabilistic logic
and use conditioned policy for learning and inference

Modelling policies as Nesy-MMS makes safe RL agents possible

Previously in our group

Program safety specification in probabilistic logic
and use conditioned policy for learning and inference

Limitation

Exact inference is too expensive for large domains
and only looks ahead one step

Modelling policies as NeSy-MMS makes safe RL agents possible

Previously in our group

Program safety specification in probabilistic logic
and use conditioned policy for learning and inference

Limitation

Exact inference is too expensive for large domains
and only looks ahead one step

Now

NeSy-MMs encode the conditioned policies
allowing safe policies in large temporal domains

Neural nets need relational probabilistic reasoning to achieve the guarantees they lack

Relational NeSy-MMs...

...are scalable and differentiable
neural + logical + probabilistic models

...are logically consistent and intervenable
while generalising better than other sequential models

...show promise for larger applications
such as controlled language generation and safe RL



Visit personal page



Read our paper